

14.09.22 Lec 6

Process Management

OS provides abstraction to user programs

→ Hardware is complex

→ Challenging to understand
how the hardware
actually works

3 fundamental abstractions that OS
provides to the user programs

→ Process

may hurt
performance
in terms of
execution

→ Virtual Memory

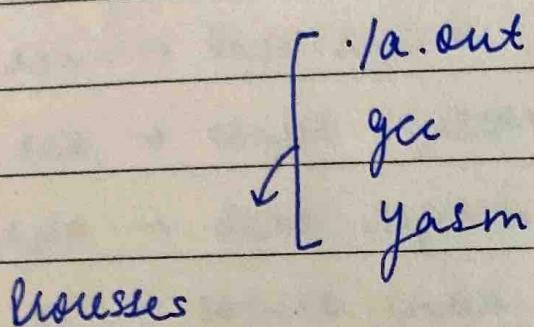
time or
additional
memory or
add. storage

→ File System

Process

when you run a program, the actual
binary has to be loaded into the main
memory. there is a "running instance" of
the program that is consuming memory,
consuming time of the processor and

making requests to I/O devices. This running instance is called process.



while they are in execution

Suppose you have a single processor,
you can run multiple processes.

→ Time-sharing

→ Job of process management

System of OS is to ensure
that no process is unfairly
hurt and they all get
fair access to the processor's
time.

1) How can it use the I/O devices?

→ Processes are > 10000 times
faster than I/O devices

→ Interrupts are needed to signal
the need for I/O

→ User programs can only raise
an interrupt & the OS has a device management

System which can directly interface with the actual I/O device

Interrupts when raised → normal execution of the program stops.

→ control of execution jumps to a specific location in the memory depending on the type of interrupt.

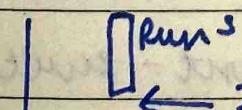
→ during bootstrap, the memory management system of OS loads specific instructions to handle each type of interrupt at different possible locations.

int #1 - ← store the specific

int #2 - interrupt service routines

int #3 - in these locations.

Time



← Interrupt

← Processor Management has to send this process back

Processor execution mode or privilege level

gui is a process

→ kernel has privilege level 0 so it can execute everything. Whereas user programs have higher values set in the privilege mode, so it cannot access directly the I/O devices

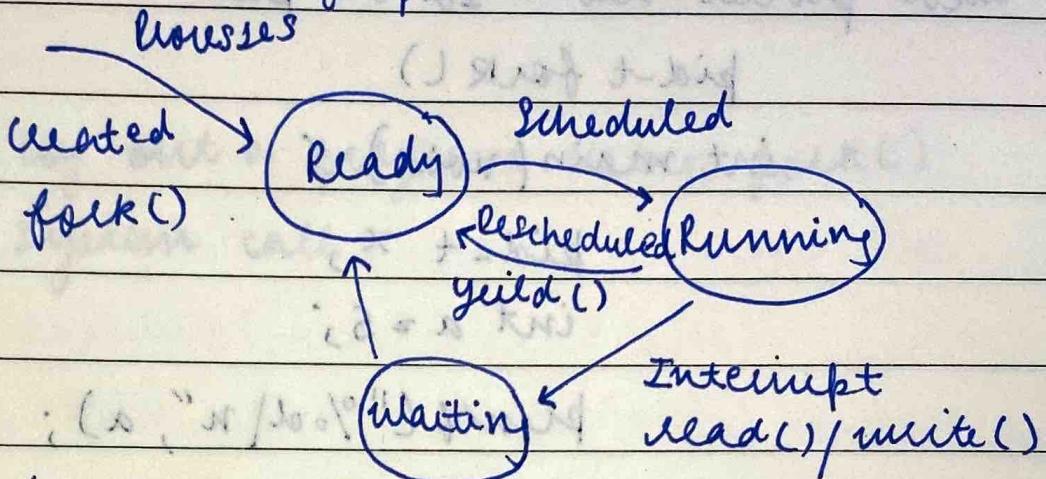
processes and their privilege levels help build isolation among different programs

All the interrupts cause a context-switch. The context switch implies that a privilege mode is being changed. That has some overhead, because you want to ensure that the execution can resume.

process-switch: → If the process management wants to stop execution of one process & start execution of another process, then it has to perform a process switch. Here the overhead is higher than in context-switch requires process management system to save the "state" of the entire process.

Data structure called (PCB) stores details of each process.

States of execution of processes



We discussed mechanisms of processes running. There are complex design choices that each OS needs to make. Called policies.

- ① Scheduling, descheduling policies decided by process scheduler.
- ② Design of the system calls → how they should be used is also a design choice.

User program → OS → hardware
System calls

In VM → every instⁿ goes thru OS (diff system calls)

Elsewhere → limited directed execution,
everything doesn't go through the OS.

The fork() system calls

for unix

↳ create a new process

Each process has a ID or pid

pid_t fork()

```
int main(void) {
```

```
    pid_t x;
```

```
    int a = 5;
```

```
    printf("%d\n", a);
```

```
    x = fork();
```

```
    if (x == 0) {
```

```
        /* child process */ printf("a=%d\n", a);
```

5
7
6
7 } -

}

```
else if (x > 0) {
```

```
    /* parent process */ printf("a=%d\n", a);
```

non
deterministic

output.

should be
avoided

```
else {
```

```
    printf("Error");
```

}

```
printf("%d", a);
```

```
return 0;
```

Race conditions → different parts of the code
all racing against each other to execute first.
avoid.

glibc

the exec system calls allows you to load
a binary into the main memory for
execution.

Exercise → Try out a program using fork()
system call.

17.09.22

Saturday

fork() creates a new address space of the child process

parallel / concurrent programs can be written

Native fork() has to copy a lot of data which is not used in practice.

→ performance optimization

copy-on-write (cow) → copy of data / code is made only when it is changed in the child / parent process.

Parallelism

Multi-core system → parallelism (MIMD)

→ because two or more "inst" are executed simultaneously.

GPU → SIMD parallelism

→ "single inst" on 2 or more data variables

can only be used in programs that have concurrency.

concurrent \Rightarrow parallelism

parallelism \Rightarrow concurrent

i.e. even if instⁿ are executed in a diff. sequence / diff. processor, the result is correct.

fork → allows to create concurrent programs and so take adv. of multi core systems threads
posix_spawn() → first calls fork() then exec() s.t. addⁿ optimizations related to system variables are possible

instead of copying, programmer wants to
the address space

control the variables after creating the new process. Can be done using wfork().
(locks on variable^s)

$x = \text{fork}();$

if ($x > 0$) {

/* Parent */ $a = a + 1$

else if ($x == 0$)

/* child */ $a = a + 2$

exec("./p.out", ...);

else { /* error */ }

```
int x[2] = {0, 0},  
y[2] = {1, 2},  
sum[2] = {0, 0};
```

```
pid_t pid;
```

```
pid = fork();
```

```
if (pid > 0) {
```

```
    sum[0] = x[0] + y[0]; printf("%d",  
}     wait();  
        sum[0]);
```

```
else if (pid == 0) {
```

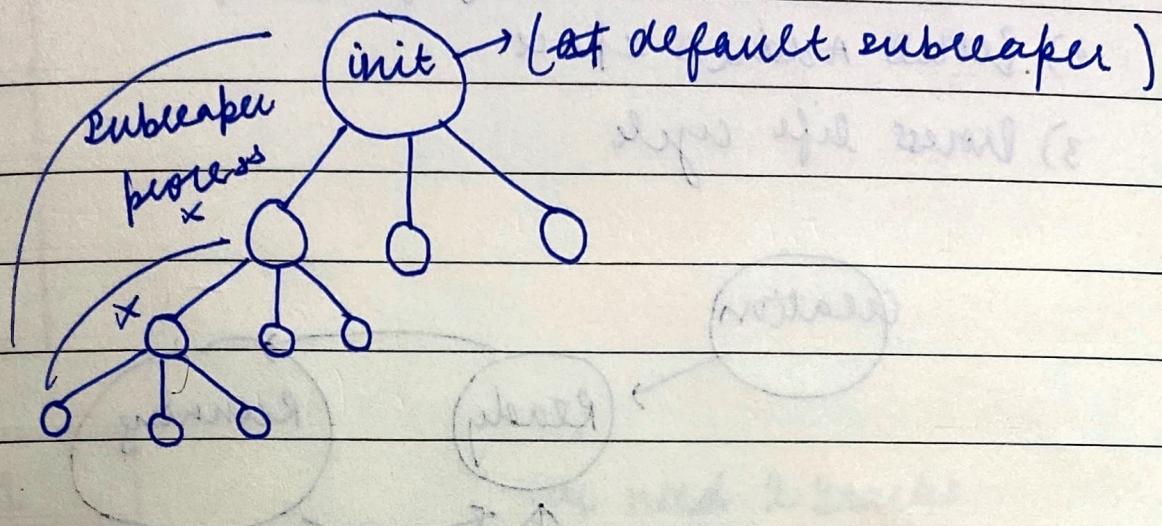
```
    sum[1] = x[1] + y[1]; printf("%d",  
}     sum[1]);
```

The `wait()` system call ensures that the parent process waits until child processes has completed its execution.

child process if its own parent finishes execution assigned a new parent.

Traditionally init process was assigned as the parent

In Linux, concept of subreaper process was created

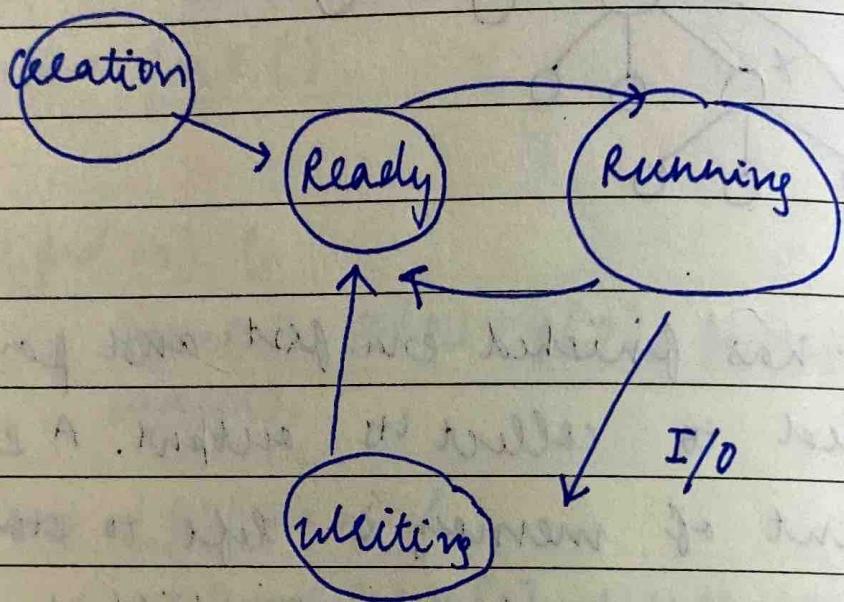


Child has finished ex. fist and parent has failed to collect its output. A small amount of memory is left to store the output and status of the child process. Such a process is called zombie process.

21.09.22

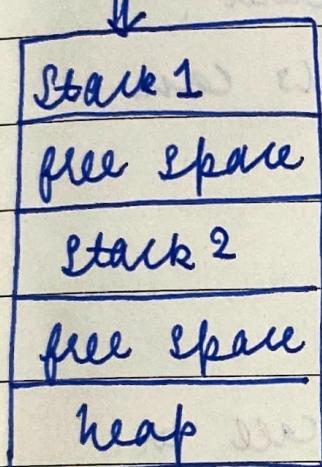
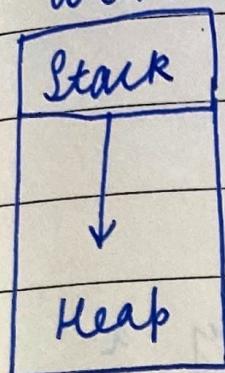
Process Management

- 1) Create process using fork
- 2) Process Address space
- 3) Process life cycle



- 01) All I/O would force the process to go to waiting state
- 02) Any type of parallelism requires substantial amount of overhead in creating a new process.

Threads help to solve disadvantages
Threads are multiple lines of execution
within a single process.



We need 2 stacks
as each thread needs
to do bookkeeping
of which f" calls
which

```
void abc();  
ijk();  
int main() {  
    xyz();  
    abc();  
    xyz();  
    thread-create(def);  
    xyz();  
}  
def();  
xyz();  
void def() {  
}  
xyz();
```

can run in parallel

Threads enable concurrent access to data structures

Effort required by us to enable such access
and manage shared data structures.

Responsibility of handling such race
conditions is on the programmer.

User-level threads → threads created by a
process internally w/o informing the kernel.
Switching execution b/w the 2 threads can
be done by the process itself.

"green threads"

↳ not an OS concept

Kernel-level threads → Use a system call to
create a thread. This system call automatically
creates a new stack & also separate PC
register values. Switching execution b/w the
threads also has to be done via system calls,
which makes it time consuming.

The `Clone()` system call creates threads
`posix-thread-create`
`posix-thread-join`

int main (void) {

posix-thread-create(abc);

1st computation of your

own thread * /
polix-thread - join(abc),

void abs() {

↳ Computation of abc

thread * /

y

```
for (i=0; i<100; i++) {  
    // code  
}
```

`sum[i] = x[i] + y[i];` ; <sup>(not
threaded)</sup>

first methods were → sequential

Concurrency when each job uses

```
void computeSum(int arr-les[], int x[], int y[],  
                int start, int end)
```

```
for( i = start; i < end ; i++ ) {
```

$$\text{sum}[i] = x[i] + y[i].$$

} without go right to street level sand

int main(void) {
 cout << "Hello world" << endl;
}

```
int x[100], y[100], sum[100]; posin_thread  
for(int i=0; i<4; i++) { *t[4]
```