

Tensorflow: internal libraries
implement concurrency

python - no support for concurrent threads

$t[i] = \text{posix-thread-create}(\text{sum}, x, y, i * 25,$
 $(i + 1) * 25);$

}

for (int i=0; i<4; i++) {

 posix-thread-join(t[i]);

}

return 0;

}

Python

→ Global Interpreter Lock (GIL)

$\text{sum} = x + y;$ → vector operation that
can be done using data
-level parallelism

When are threads most useful?

→ Browsers w multiple tabs

→ Servers with many requests

Linux keeps track of type of machine
during install and does scheduling of
processes and threads accordingly.

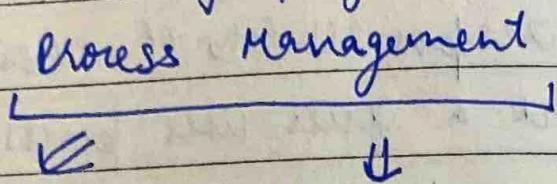
Linux assigns a new process ID for every thread. The responsibility of ensuring that an I/O doesn't block a full user process is on the scheduler, which checks if other threads of the same user are available

new ①

newselect ②

26.09.22

Assembly programming



- ① which system ② concepts of
calls to use ? process/thread

③ Internal organisation

within the kernel

④ roses
scheduler

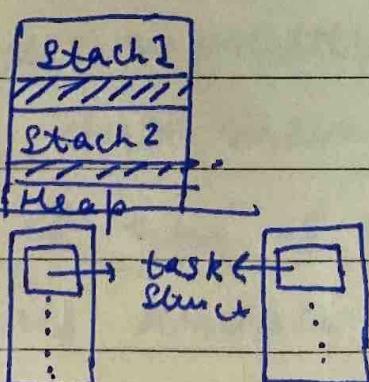
③ Inter process communication & Synchronization

Address Space → Memory area where we store the process, data structures.

related

User Space

Kernel Space



(for each
thread)

Internal data structures related to each process kept by the kernel. Kept in the process address space within the kernel space.



task → anything that can be scheduled
thread/process

In linux, process control block is called
task-struct.



You can enter into (execute code from)
kernel space in two ways :

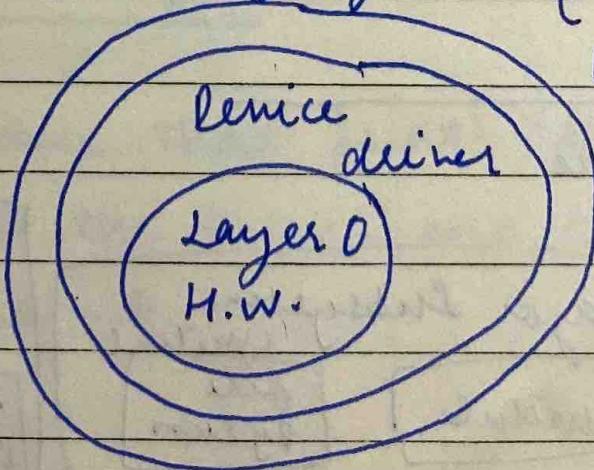
- 1) Process context → Enter into it
using system calls.
- 2) Interrupt context → Enter into
using hardware signals

How does the kernel handle requests?

→ Linux is a monolithic kernel

↳ any fn in the
kernel can call
any other fn in
the kernel

→ Layered operating system (UNIX system V,
Berkeley UNIX)



→ Microkernel → More file system & device drivers to the user space. The OS is secure from crashes due to device driver bugs. Very reliable; used in embedded & safety-critical systems.

GNU Hurd → Attempt to create a microkernel for desktop systems

Linux Kernel Modules → Parts of the kernel that are not needed to run the basic codes of processes. Modules can be loaded when the ^{kernel} OS is running. They can also be removed if it is not in use. Once loaded, the modules can access any Linux function.

Kernel space
(common)

all ^{to} processes

Major Subsystems

Schedule

Virtual
file
system

other
components

Arch
specific

M V

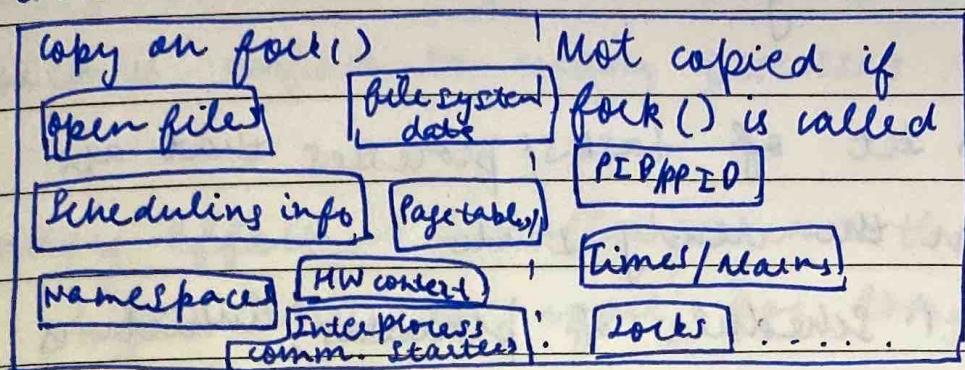
Device drivers

Network

Character

Storage

task_struct



Circular linked lists to keep track of task_structs.

There is a "current" variable which keeps track of the current task_struct. This current variable is stored in a register. Whenever the kernel wants to make a process switch, the current variable has to be modified.

Terminating a process for a kernel is the same as destroying the task_struct.

A kernel module that prints "Hello world". Load it into the kernel and see whether message comes up.

Parent process ID
↓
process + P

28.09.22 Wed
/Process Scheduling
Task

Input: A set of tasks/processes that are in the ready state

Output: A schedule of processes, which is the same as selecting a process to send it to the running state.

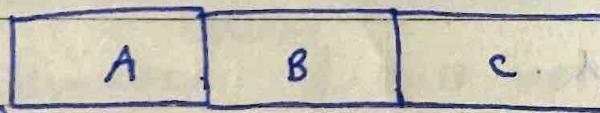
First come first serve (FCFS) → Non-preemptive

A → 15 time units, arrived at 0

B → 20 time units, " 1

C → 10 time units, " 2

Burst time/execution time



Gantt chart

Mean turnaround time = Avg time taken to complete the task

Completion time -
Arrival time -
Burst time
System time

Since it's beginning

$$\text{Turnaround (A)} = 15$$

$$T_{\text{turnaround}} (B) = 35 - 1 = 34$$

$$T_{\text{turnaround}} (C) = 45 - 2 = 43$$

$$\text{Turnaround (Avg)} = \frac{15 + 34 + 43}{3}$$

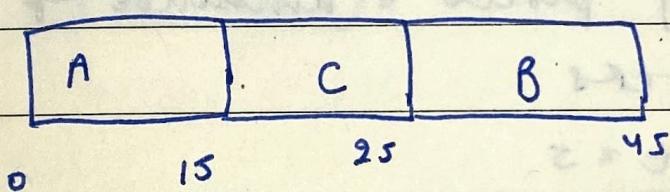
$$WT = TAT - BT$$

↳ Pre-emptive / Non pre-emptive
time quantum,
desponsiveness,
Priority

Average start time of FCFS can be large,
which is unfair to many processes.

Convo effect → smaller processes that come
on might have to wait for long time
for a long process to end finish.

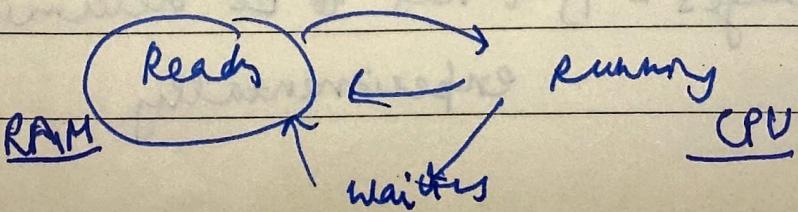
Shortest job first criteria :
schedule at each step the job that will
take the shortest time to finish.
burst time
Non-preemptive



Starvation

A process may not get access to the CPU
for infinite time.

So far, the assumption is that the process
cannot be moved back ^{to} from the ready state
from the running state i.e. processes are
non-preemptive in nature

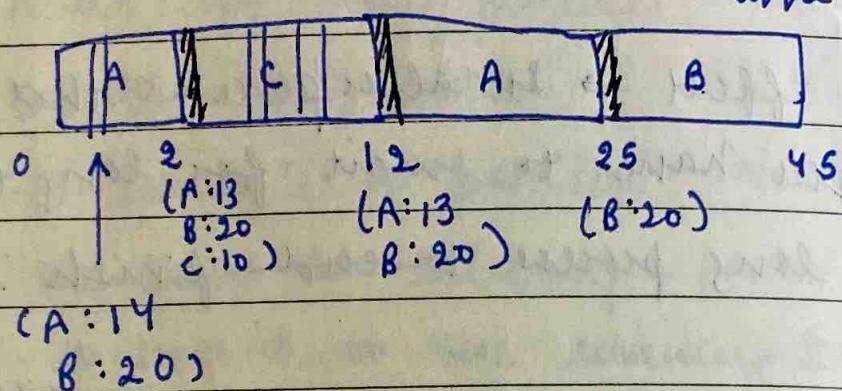


preemptive version of SJF \rightarrow [Shortest Job

"Scheduling"

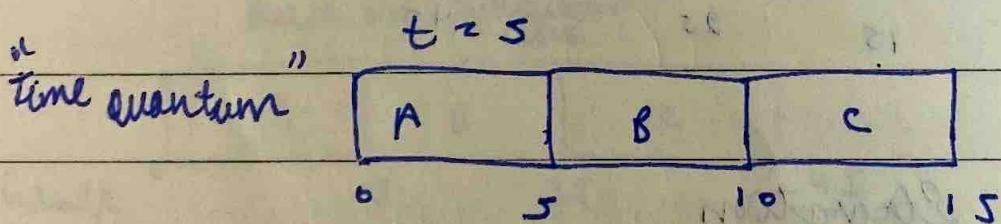
remaining \downarrow first]

time



Preemptive

Round-Robin Scheduling: Divide some unit of time called time quantum t . Allocate each ready process t amount of time based on FCFS.



Advantages - 1) No process is kept waiting for too long.

2) No need to have accurate estimates of process execution type.

Disadvantages - 1) t has to be determined experimentally

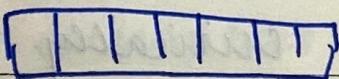
small burst time

2) If t is large, then interactive jobs are not satisfied. If t is small, then there are too many process switches.

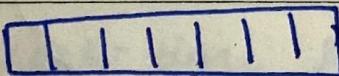
3) If t is known, then a process might be written in such a way that it uses I/O only just before its t finishes.

Multi-level Feedback queue

Q1



Q5



- ① Give higher priority to a process that is at a higher queue \Rightarrow if P_1 is at a higher queue than P_2 , then schedule P_1 .
- ② Higher queue may a smaller time quantum than a lower queue

③

Let every process start from Q1. If it uses the entire time quantum, then move it to lower queue for its next execution. If it does not use the entire time quantum, then keep it in the same queue.

④ Starvation-Free: After some S amt of time, send all the lower-queue processes to the highest queue.

Scheduler performance is critically dependent on the heuristic values
(time quanta and S)

The schedulers discussed today are fine for single processes. For multiprocessor scheduler, there needs to be proper handling of additional race conditions.

3.10.22, Monday

- 1) what an OS does? → Principles
 - 2) how does it perform its actual actions?
 - Latest version of the linux kernel
-

Process Scheduler

↳ 2 types

- 1) Normal
 - 2) Real-time scheduler
- ask for some specific guarantees
Latency guarantees

Soft real-time schedulers → (playing video)

Hard real-time schedulers → (Industries)
↓
cars, etc

Linux kernel cannot

satisfy these constraints

How can the Linux kernel handle both soft-real time constraints & normal jobs together?

- Every process must belong to a scheduling class.
- Scheduling class defines the scheduling policy that will be used.

→ Each class assumes that processes/threads are preemptive

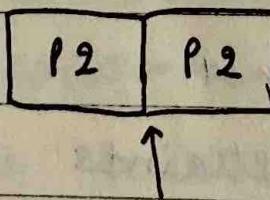
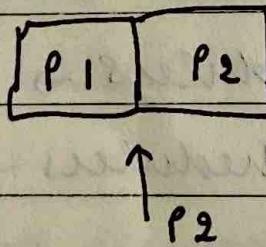
- 1) SCHED-FIFO → uses FIFO scheduling with no time slice;
if higher priority job in this class comes, then this job can be moved back to the steady state

Priority:

50 (P1)

60 (P2)

(priority values from 0 to 99)



when P2

ends, P1 starts

- 2) SCHED-RR (soft real-time) → very similar to SCHED-FIFO, but w a timeslice of 10 ms after which a process is preempted.

3) SCHED-NORMAL (Normal user processes) →

Priority values of 0; but "nice" values can be used to change priority.

Nice values → Higher values imply lower priority

-20 to +19

default is 0

A user can increase the nice value to indicate that the process can be slower than usual.

only root users can set nice values to (-)ive
↳ not ordinary users

2.6.23 of the kernel → Completely Fair Scheduler (CFS)

Provide equal amt. of access to the processor to each task

160 ms

50 ms 50 ms

10 ms . . .

10 ms

Too many process switches are possible if a fully fair scheduler is used. To avoid it, a floor of 1 ms is kept below

* imp. part

which time slice is not divided.

"Amount of access" is the time spent at the processor (stored in a variable called vruntime).

At each point of time, choose the process w the lowest vruntime , If kept in a linked list / array, it can take $O(n)$ time.

Red black trees to store vruntimes.

$$O(\log n) \times 2$$

which reduces the time complexity to $O(\log n)$.

There is a table of weights based on the nice values, where higher nice values have higher weights.

$$[\text{vruntime } \leftarrow \text{actual - processor-time}]$$

$$\star \text{ weight [nice-value]}$$

System call called " nice" which can be used to change the nice value of a process.