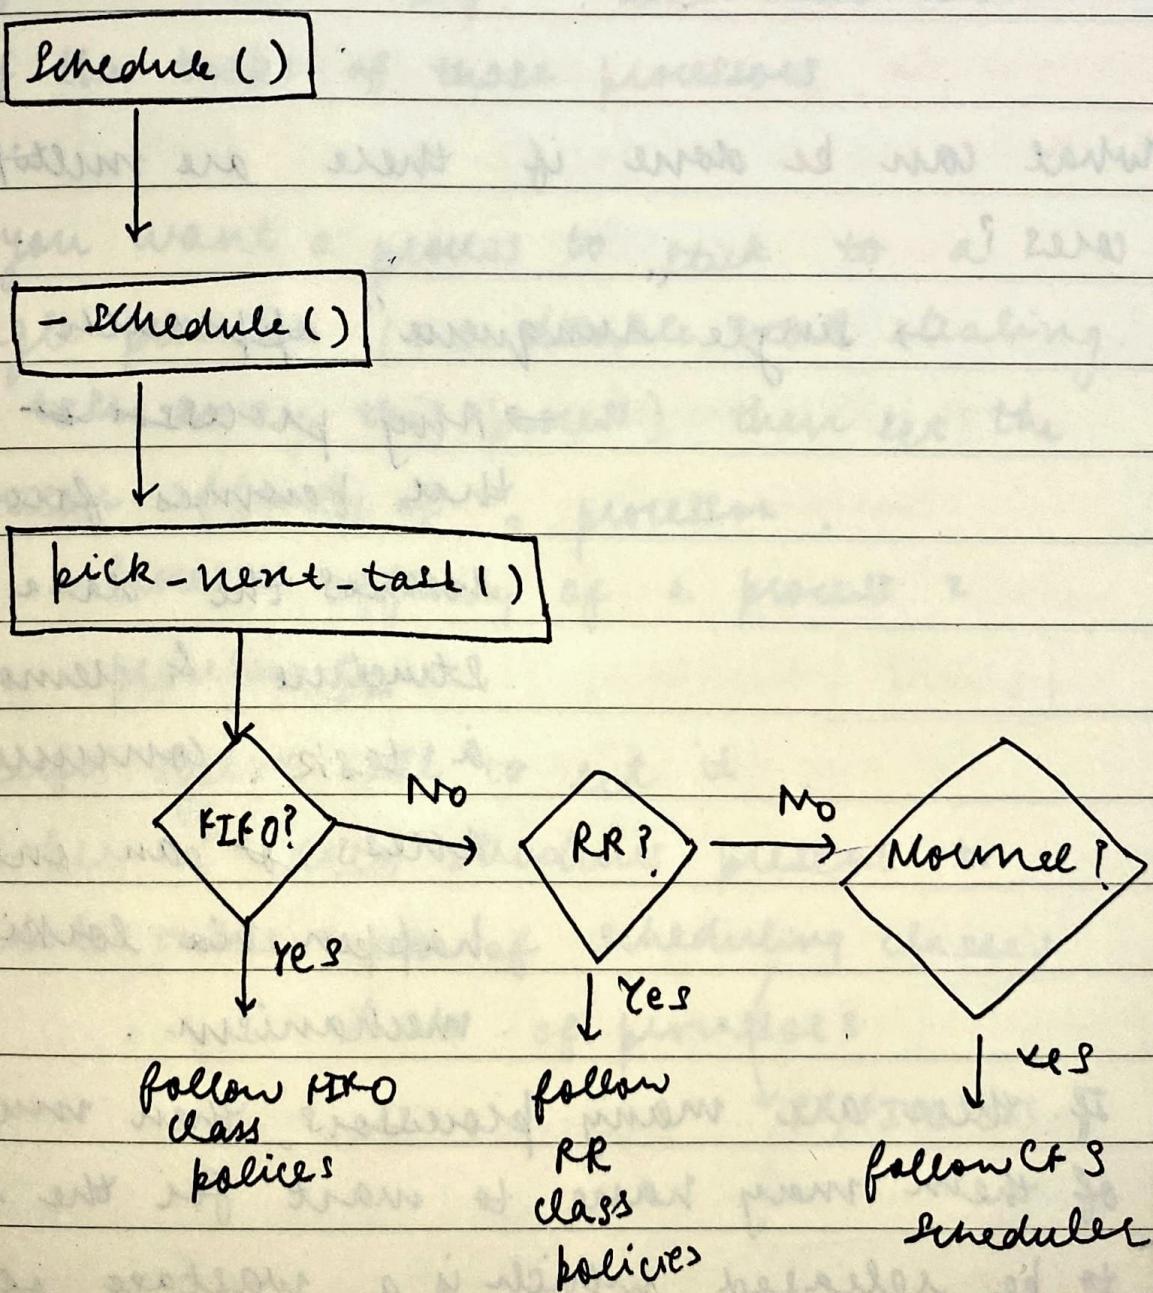


How exactly are processes scheduled across classes?

There is a priority among classes. If a process w higher "class priority" is present, that process will always be selected.



FIFO, RR classes have their own linked lists; their processes are not kept in the red-black tree.

CFS eliminates the need to separately classify processes based on interactive & non-interactive.

what can be done if there are multiple cores?

→ Single "runqueue" approach  
→ Any processor → that becomes free  
accesses the data  
structure & removes  
a task. concurrent  
access → can only  
happen via locking  
mechanism.

If there are many processors, then multiple of them may have to wait for the lock to be released, which is a waste of processes time.

Cache effect

→ Individual "uniqueness" approach →

Each processor has its own set of processes; and the tasks are distributed among the processors.

An ideal processor would check the uniqueness of other processors, if found non-empty, it will take some of the tasks of those processors.

If you want a process to stick to a single processor (i.e. disable work stealing to take away this process), then set the processor affinity of a processor.

Processor affinity of a process = processor ~~#~~

Specific system call to set it.

Number of data structures present to schedule tasks = No. of scheduling classes  
X No. of processors

PR, FIFO etc.

10.08.2022, Monday

1) Assembly

2) Process Management

→ Process Data Structures

→ Process System Calls

→ Threads

→ Process Scheduling

3) Structure of kernel → layered kernel, microkernel

4) How do system calls work?

5) How does the bootup process work?

System calls

→ A mechanism of a user program

communicating with the OS.

`printf` → library function defined by

the library glibc.

[uses embedded assembly]

`wlfe` → library function with the same name as

-----  
|-----  
|----- system call

`wlfe` (System call) → System call defined by the Kernel

↓ kernel memory  
Space starts here

(no access to library functions)

{ copy-to-user to copy the data <sup>to</sup> from the user's memory space <sup>to</sup> kernel's memory space.

copy-from-user to copy the data <sup>to</sup> from the kernel's memory space <sup>from</sup> user's memory space

→ defined by linux kernel itself

are functions in the kernel

write

⇒

System call number

this is done in

process context, so  
multiple processes

can make similar  
requests at the

same time

Find correct system

call from the

System call table

↓  
check the arguments  
whether they are  
valid

processing of requests  
by making requests to device drivers or kernel data structures.  
↓  
copy-from-user

copy-to-user



copy return value to  
user register



return control to  
user program

## Boot Loader

1) Processors start in 16-bit mode (real mode)

→ No proper memory  
management

→ BIOS (Basic I/O System)

→ checks the hardware

present; on error it stops  
the boot process.

processor, RAM, keyboard

BIOS stored on ROM; needs to load  
data from disk.

Small disk drivers are present in BIOS to  
access the disk.

→ first sector of the disk needs to be in a specific format if it has the OS loaded into it.

Linux has a bootloader called GRUB.

→ Boot loader takes control from the BIOS.

→ Reads each of the disks to find out the OSes present.

→ If Linux kernel is chosen, then it loads the kernel's image into the main memory.

→ Sets up the data structures required to enable memory management unit & then sets a flag in the control register CR to move to 32-bit mode (protected mode).

Linux kernel has different levels of service called "init levels".

All GUI + other processes allowed → init level 5  
Everything except GUI → init level 3

Single user → init level 1  
GUI and network → init level 2

The init process is the first process created; the rest of the processes are spawned by using fork() and exec() one-by-one.

The command "init 0" shut down the machine.

Daemon process: A type of process that runs in background and provides different user services (eg:→ networkd, sshd).  
Bash shell not daemon process.

boot process

BIOS: Basic I/O System

↳ 16-bit operations

→ Hardware checks

1) Booting process was slowed down by moving to 32-bit/64-bit mode.

2) Firmware can be upgraded by downloading (ROM) and flashing the newer images

↳ introduces a security risk by allowing malware to be installed

The idea of trusted computing came in where an OS had to "sign" to be allowed access into the firmware.

only if the signature (digital key) is present, then access is allowed. This makes it difficult for a new OS to be booted

Master boot record (MBR)

UEFI → Unified Extensible Firmware Interface

Booting from multiple partitions of same disk is possible.

hardware check for more number of hardware is possible.

12 Q.S.

8 3 mark Q.S.

2 5 mark Q.S.

2 8 marks

50

90 mins

1) a) when is FCFS scheduling not good for processes in terms of turnaround time?

b) what alternate scheduling can resolve this?

create new process & load it  
process / threads

Robert love : Linux Systems Programs  
↳ talking to the kernel

31.10.22

Assembly programming → apart from arithmetic, logical ops, branches & loops, concept of memory alignment

Process Management → lifecycle, data structures, threads, scheduling

Structure of OS

Bootstrap process

Process Management

How can processes coordinate among each other? → Since processes are isolated, then why is process communication needed?

⇒ Inter-process communication (IPC) is needed to coordinate access to shared resources

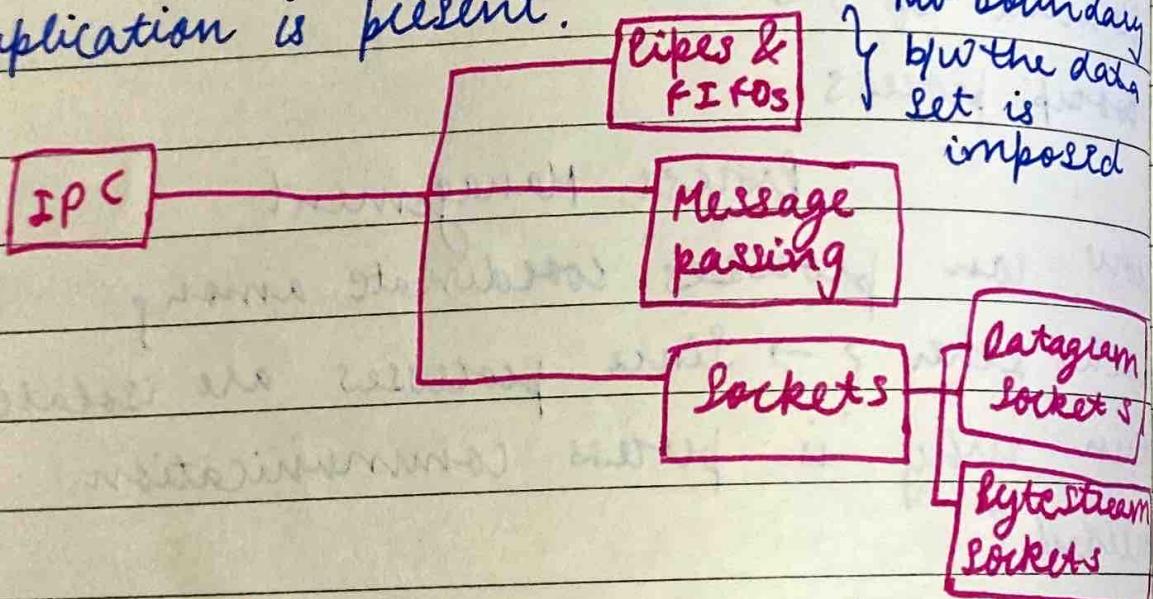
All data communication existing today uses IPC.

POSIX is the common standard used by all UNIX-based OS to provide process communication mechanisms

System V standard coming from AT&T's version of UNIX.

⇒ Also provided mechanisms of IPC.

Most UNIX systems have support for IPC from 2 different libraries, we cover only POSIX libraries, whenever duplication is present.



Communication through files is considered to degrade performance.

① Most of the IPC mechanisms in UNIX use the standard file-handling machinery. The advantage of using it is that there is no need to build familiarity with a new set of system calls.

① Shared memory → If you need to use many read's & writes, then the file handling mech can reduce performance due to context switches. If you use shared memory, then reading & writing is similar to reading & writing from an array. But consistency has to be ensured by the user processes through use of proper locks.

pipe ("Anonymous") doesn't have location, kept in kernel memory space  
 → Virtual or pseudo-file which can be used for IPC returns 2 functions  
 $fd = \text{pipe}(\text{fd}[2]);$

$\text{int fd}[2];$

$\text{write}(fd[0]);$

$\text{read}(fd[1]);$

$\text{pid} = \text{fork}();$

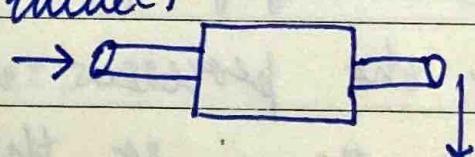
$\text{if} (\text{pid} == 0)$  { /\* child process \*/  
 $\text{char}^* \text{buf} = \text{"abc"};$   
 $\text{write}(fd[0], \text{buf});$

else {

$\text{char}^* \text{buf};$

$\text{buf} = \text{malloc}(100);$

$\text{read} = (\text{fd}[1], \text{buf});$  }



read()

- 1) Disadvantage of pipe is that there is no way of communicating across processes that are not related.
- 2) Each process can only read or write but not both. Communication is unidirectional. Using a single pipe bidirectional communication would require 2 different pipes.

FIFO or named pipe can be used to communicate across processes that are not related.

fifo ("abc");

Both the processes would need to use the same name so that the kernel can identify that they want to communicate with each other. After opening using the system call, read/write/close can be used as usual. Other facilities are similar to the original pipes.

pipes do not guarantee that they would read/write the entire data at a time.

ret-val = write(fd, buf, buf-en)

512/1024, 4096

/\* ret-val indicates how many bytes have been written \*/

total-char = 0;

while (total-char < 1024) {

    ret-val = write(fd, buf[totalchar],  
                  1024-totalchar);

    total-char += ret-val;

}

close(fd);

Read to a pipe is blocking in nature.

while (tot-char < 1024) {

    ret-val = read(fd, buf[totchar],  
                  1024-tot-char);

    tot-char += ret-val;

}